

Change Summary

- Project: codexa
- Checkout: isolated codexa task worktree
- Branch: codex/general/codexa-019f985f-f2aa-7853-a8ec-78ad4943ea5c
- Base: main
- Primary commit: 76688b4
- Subject: feat(workflow): harden selective Codexa execution

Changed Files

76688b4 feat(workflow): harden selective Codexa execution

```
benchmarks/agent-ab-selective-v2/README.md      | 53 +++
.../config/adaptive-auto-bounded.mcp.json       | 8 +
.../config/adaptive-auto-bounded.md             | 18 +
.../config/adaptive-auto-legacy.mcp.json        | 8 +
.../config/adaptive-auto-legacy.md              | 8 +
.../config/full-detailed-legacy.mcp.json        | 8 +
.../config/full-detailed-legacy.md              | 8 +
benchmarks/agent-ab-selective-v2/experiment.json | 92 +++++
.../environment/Dockerfile                      | 21 ++
.../environment/project/.gitignore              | 3 +
.../environment/project/README.md               | 10 +
.../environment/project/src/init.py             | 1 +
.../environment/project/src/targets.py          | 11 +
.../environment/project/tests/testtargets.py    | 14 +
.../start-codexa-mcp-adaptive-auto-bounded.sh   | 41 +++
.../start-codexa-mcp-adaptive-auto-legacy.sh    | 41 +++
.../start-codexa-mcp-full-detailed-legacy.sh    | 41 +++
.../tasks/path-target-normalization/instruction.md | 14 +
.../solution/solution.patch                    | 36 ++
.../path-target-normalization/solution/solve.sh  | 4 +
.../tasks/path-target-normalization/task.toml    | 32 ++
.../path-target-normalization/tests/Dockerfile  | 16 +
.../tests/baseline/.gitignore                  | 3 +
.../tests/baseline/README.md                   | 10 +
.../tests/baseline/src/init.py                 | 1 +
.../tests/baseline/src/targets.py              | 11 +
.../tests/baseline/tests/testtargets.py        | 14 +
.../tests/candidaterunner.py                   | 67 ++++
.../tests/publictestrunner.py                  | 53 +++
.../tasks/path-target-normalization/tests/test.sh | 7 +
.../path-target-normalization/tests/verify.py   | 398 ++++++
docs/architecture/codexa-context-server.md      | 13 +-
docs/architecture/session-memory.md             | 56 +-
docs/guides/agent-ab.md                        | 110 +++++-
scripts/agent-ab-analysis.mjs                   | 322 ++++++
scripts/agent-ab.mjs                           | 118 +++++-
src/language.ts                                | 3 +-
src/mcp/compaction.ts                          | 8 +-
src/mcp/decision-kernel.ts                     | 31 +-
src/mcp/decision-policy.ts                     | 45 +++
```

src/mcp/review-coverage-kernel.ts	44 +++
src/post-edit-outcomes.ts	47 ++-
src/post-edit-review-coverage.ts	164 ++++++++
src/prove.ts	12 +
src/query-data.ts	1 +
src/query/post-edit.ts	42 ++-
src/query/post-edit/decision.ts	21 +-
src/query/post-edit/lifecycle.ts	5 +
src/query/post-edit/next-actions.ts	9 +-
src/query/post-edit/support.ts	22 ++
src/task-lifecycle.ts	3 +
src/types/query-data.ts	10 +-
src/types/verification.ts	26 ++
tests/agent-ab-report.test.ts	36 ++
tests/agent-ab-v2-identity.test.ts	251 ++++++++
tests/agent-ab-v2.test.ts	27 +-
...r-01-ranks-transitive-import-hubs-above.test.ts	41 +++
...exer-06-keeps-planned-post-edit-reviews.test.ts	18 +
tests/language.test.ts	29 ++
tests/mcp-decision-kernel.test.ts	84 ++++
tests/mcp-result-budget.test.ts	14 +
tests/post-edit-review-coverage-validation.test.ts	263 ++++++++
tests/post-edit-review-coverage.test.ts	347 ++++++++
tests/task-lifecycle-ordering.test.ts	64 ++++

64 files changed, 3216 insertions(+), 122 deletions(-)

```

create mode 100644 benchmarks/agent-ab-selective-v2/README.md
create mode 100644 benchmarks/agent-ab-selective-v2/config/adaptive-auto-bounded.mcp.json
create mode 100644 benchmarks/agent-ab-selective-v2/config/adaptive-auto-bounded.md
create mode 100644 benchmarks/agent-ab-selective-v2/config/adaptive-auto-legacy.mcp.json
create mode 100644 benchmarks/agent-ab-selective-v2/config/adaptive-auto-legacy.md
create mode 100644 benchmarks/agent-ab-selective-v2/config/full-detailed-legacy.mcp.json
create mode 100644 benchmarks/agent-ab-selective-v2/config/full-detailed-legacy.md
create mode 100644 benchmarks/agent-ab-selective-v2/experiment.json
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/Dockerfile
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/project/.gitignore
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/project/README.md
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/project/src/init.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/project/src/targets.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/project/tests/testtargets.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/start-codexa-mcp-adaptive-auto-bounded.sh
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/start-codexa-mcp-adaptive-auto-legacy.sh
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/environment/start-codexa-mcp-full-detailed-legacy.sh
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-

```

```

normalization/instruction.md
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/solution/solution.patch
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/solution/solve.sh
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/task.toml
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/Dockerfile
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/.gitignore
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/README.md
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/src/init.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/src/targets.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/tests/testtargets.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/candidaterunner.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/publictestrunner.py
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/test.sh
create mode 100644 benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/verify.py
create mode 100644 src/mcp/review-coverage-kernel.ts
create mode 100644 src/post-edit-review-coverage.ts
create mode 100644 tests/language.test.ts
create mode 100644 tests/post-edit-review-coverage-validation.test.ts
create mode 100644 tests/post-edit-review-coverage.test.ts
create mode 100644 tests/task-lifecycle-ordering.test.ts

```

Verification

```

- git diff --check: passed
- rtk npm run check: passed
- rtk npm run smoke:package: passed
- rtk proxy node scripts/agent-ab.mjs validate --config benchmarks/agent-ab-
selective-v2/experiment.json: passed
- (cd benchmarks/agent-ab-selective-v2/tasks/path-target-normalization/environment/project
&& PYTHONDONTWRITEBYTECODE=1 python3 -m unittest discover -s tests -v): passed
- SOLUTIONCHECKDIR="$(mktemp -d /tmp/codexa-solution-check.XXXXXX)"; mkdir -p
"$SOLUTIONCHECKDIR/src"; cp benchmarks/agent-ab-selective-v2/tasks/path-target-
normalization/tests/baseline/src/targets.py "$SOLUTIONCHECKDIR/src/targets.py"; git -C
"$SOLUTIONCHECKDIR" apply --check --unidiff-zero "$PWD/benchmarks/agent-ab-
selective-v2/tasks/path-target-normalization/solution/solution.patch"; git -C
"$SOLUTIONCHECKDIR" apply --unidiff-zero "$PWD/benchmarks/agent-ab-selective-v2/tasks/path-
target-normalization/solution/solution.patch"; PYTHONPATH="$SOLUTIONCHECKDIR"
PYTHONDONTWRITEBYTECODE=1 python3 -m unittest discover -s benchmarks/agent-ab-
selective-v2/tasks/path-target-normalization/environment/project/tests -v: passed

```

- Codexa post-edit-review: generated as a local verification artifact
- Codexa verdict replan (none) accepted over block: The saved machine snapshot intentionally covered the core authority paths rather than the final four-workstream diff; all six invariants have explicit evidence, the full deterministic gate passes, and the budgeted protocol and authority reviews are clean after the substitution-path fix.
- git diff --cached --check: passed
- staged safety scan: passed

Notes

Fixes 119

Implements four bounded improvements:

- fail-closed post-edit coverage bound to task, snapshot, candidate order, and analyzed targets;
- a public no-spend schema-v2 four-arm conformance pack pinned to Codexa 0.17.0;
- recovery-only session-memory guidance;
- truthful file-only C, C/C++, Ruby, and PHP lanes.

Risk-budgeted review: the protocol specialist was clean; the authority specialist found a target-substitution laundering path, the validator and all-surfaces regression were corrected, and the affected rereview is clean. No new dependency, source-writing MCP capability, paid model call, or configuration toggle is introduced. Release deployment remains on the normal Release Please lane.